UNIVERSITY OF SCIENCE - VNU, HCM CITY

INTRODUCTION TO BIGDATA

LAB1 - A GENTLE INTRODUCTION TO HADOOP

Students:
Vo Trung Hoang Hung
ID: 21120011

Tran Ky Thanh
ID: 21120556

Huynh Cong Triet
ID: 21120577

Hoang The Trung
ID: 21120583

Teachers:
Nguyen Ngoc Thao
Do Trong Le
Bui Huynh Trung Nam

Contents

1 Team's Result

Section	Percent	Notes
1. Setting up SNC	100%	
2. Paper Reading	100%	
3. Running Word Count	100%	
4. Bonus	100%	

2 Answer

2.1 Setting up Single-node Hadoop Cluster

2.1.1 21120011:

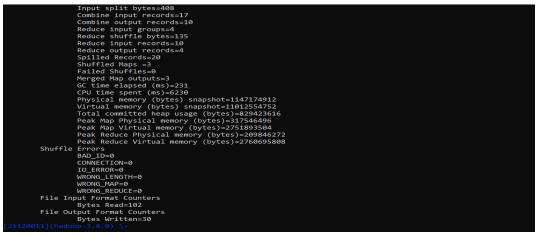

(a) hadoop-mapreduce-examples-3.4.0.jar

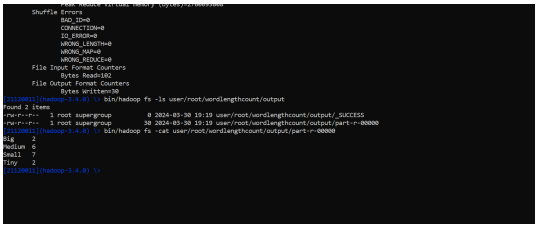

(b) List all files in output folders

Figure 1: Standalone Operation

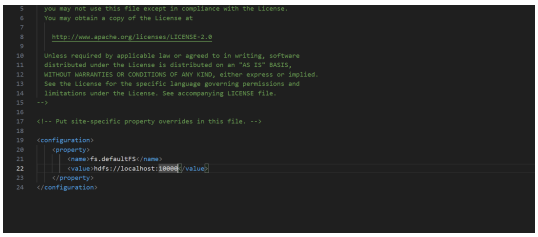

(a) core-site.xml

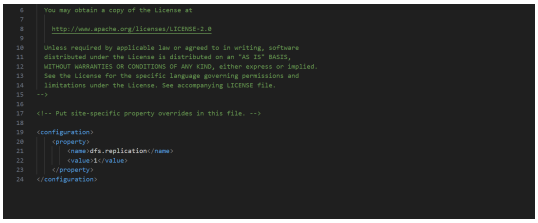

(b) hdfs-site.xml

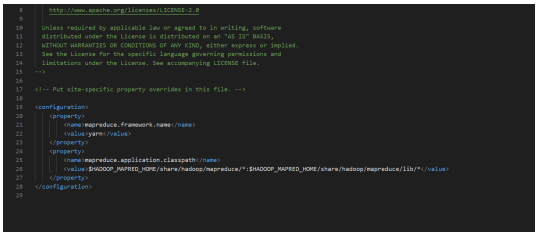

(c) mapred-site.xml

(d) yarn-site.xml

Figure 2: Hadoop configuration

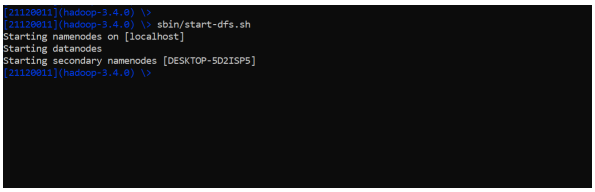

(a) Start Dfs

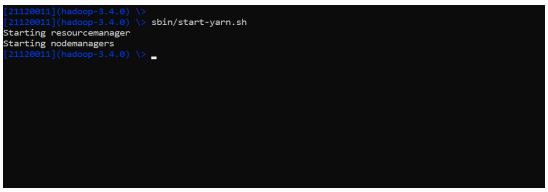

(b) Start Yarn

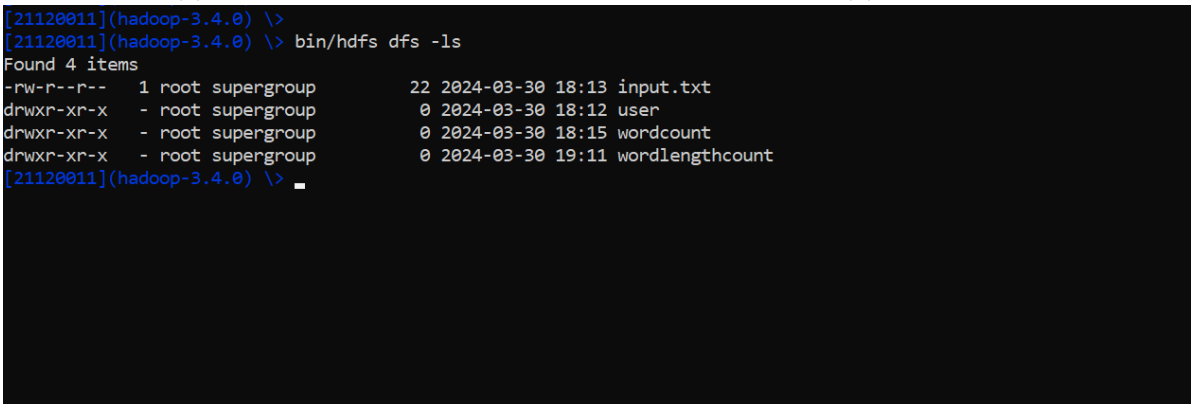

(c) Create Folders

(d) JPS

Figure 3: Pseudo-Distributed Operation

2.1.2 21120556:

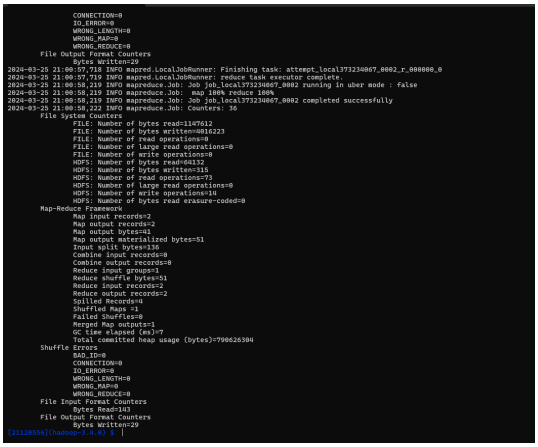

(a) hadoop-mapreduce-examples-3.4.0.jar

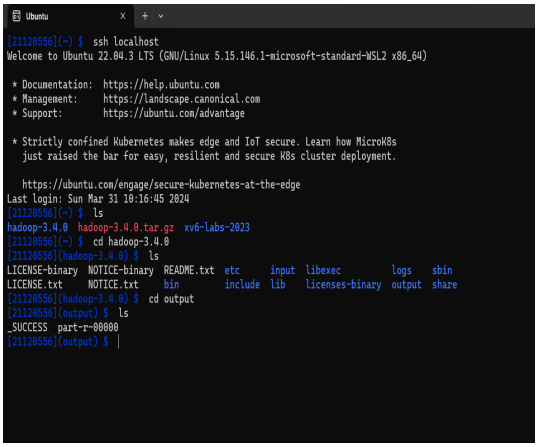

(b) List all files in output folders

Figure 4: Standalone Operation

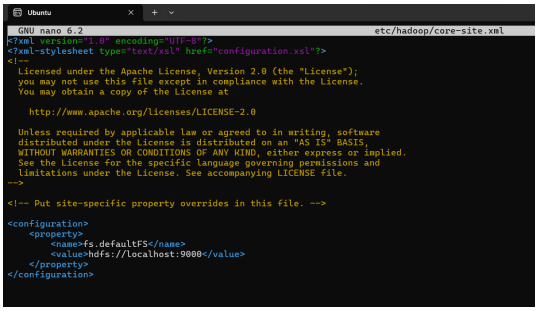

(a) core-site.xml

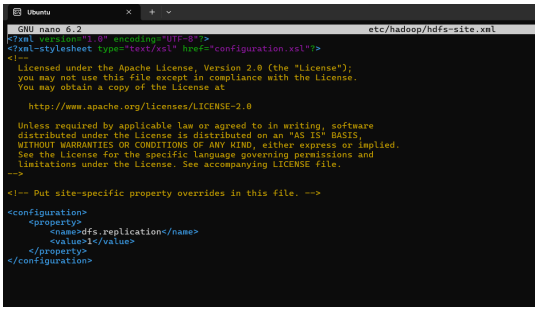

(b) hdfs-site.xml

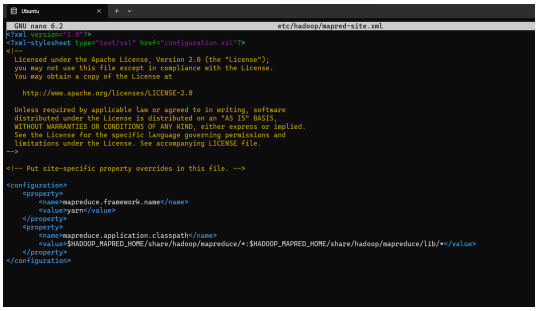

(c) mapred-site.xml

(d) yarn-site.xml

Figure 5: Hadoop configuration

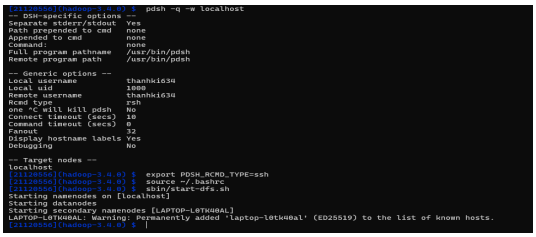

(a) Start Dfs

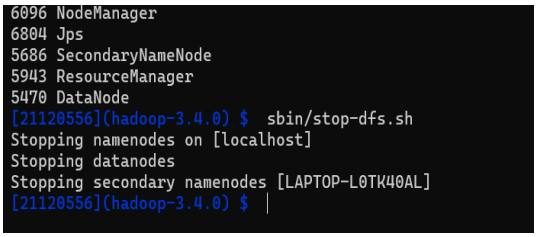

(b) Stop Dfs

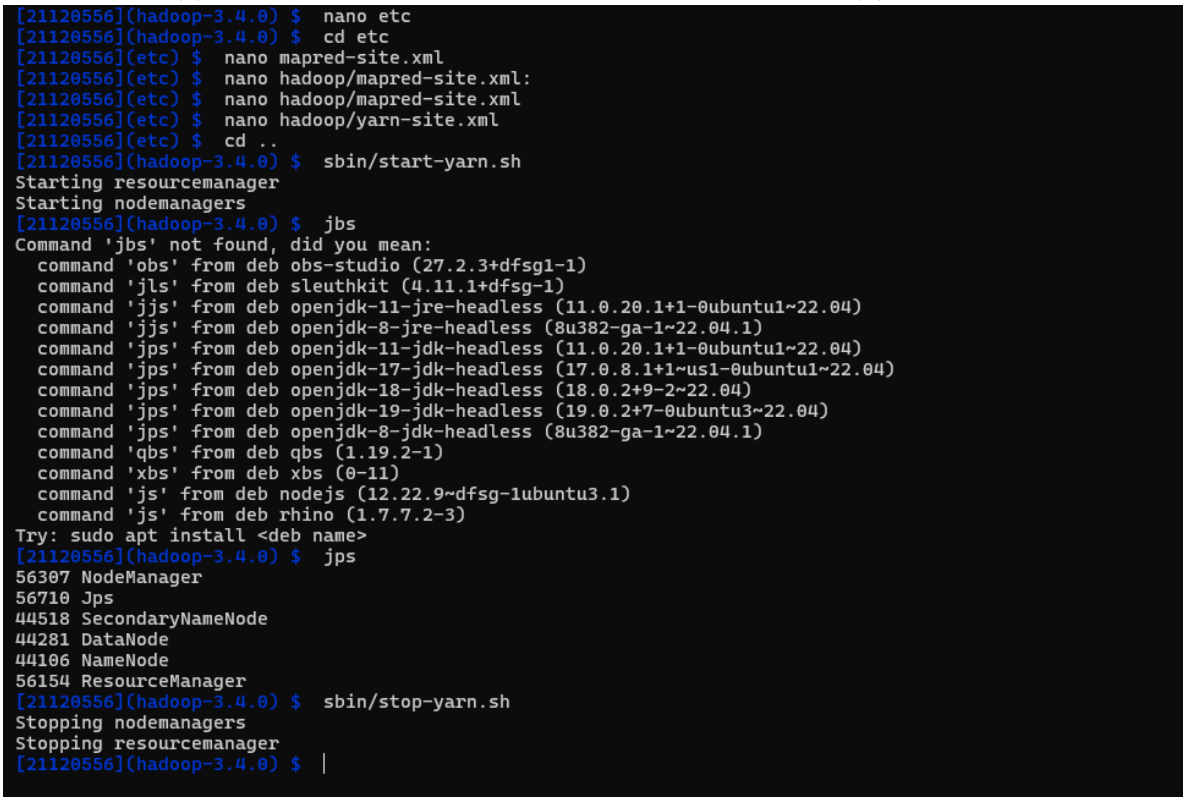

(c) JPS and Start/Stop Yarn

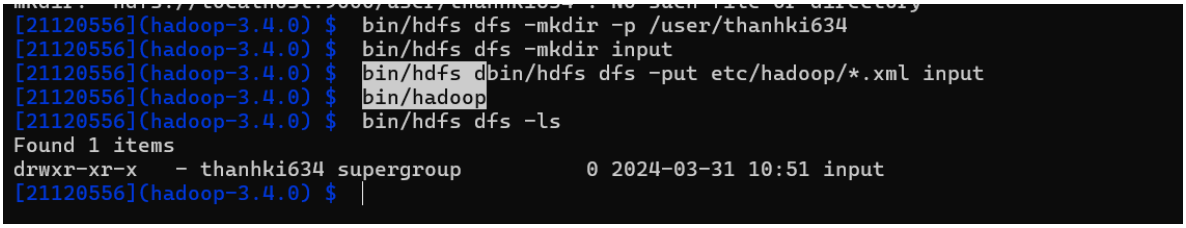

(d) Create Folder

Figure 6: Pseudo-Distributed Operation

2.1.3 21120577:

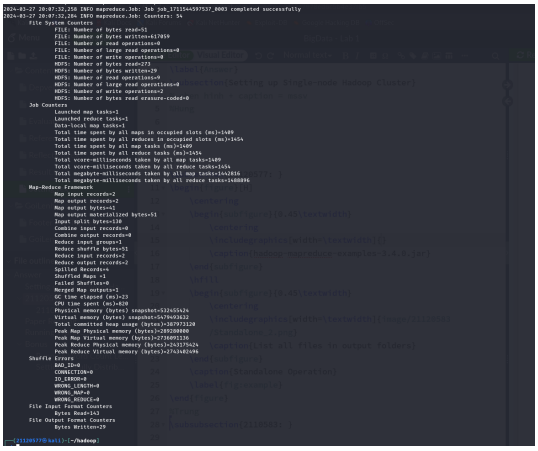

(a) hadoop-mapreduce-examples-3.4.0.jar

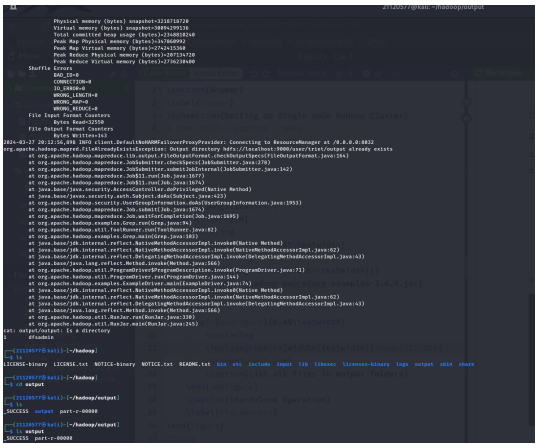

(b) List all files in output folders

Figure 7: Standalone Operation

(a) core-site.xml

(b) hdfs-site.xml

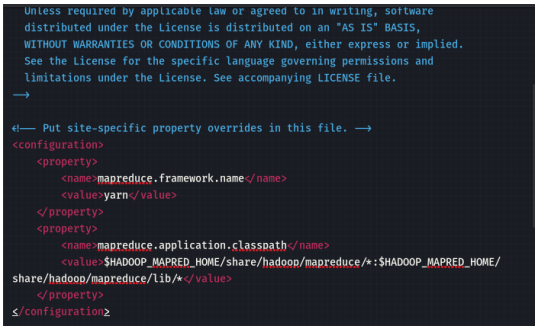

(c) mapred-site.xml

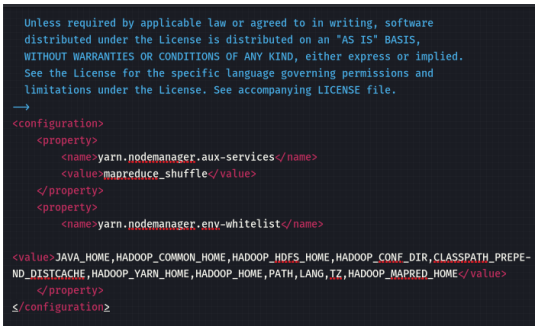

(d) yarn-site.xml

Figure 8: Hadoop configuration

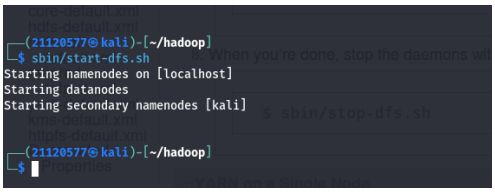

(a) Start Dfs

(b) Start Yarn

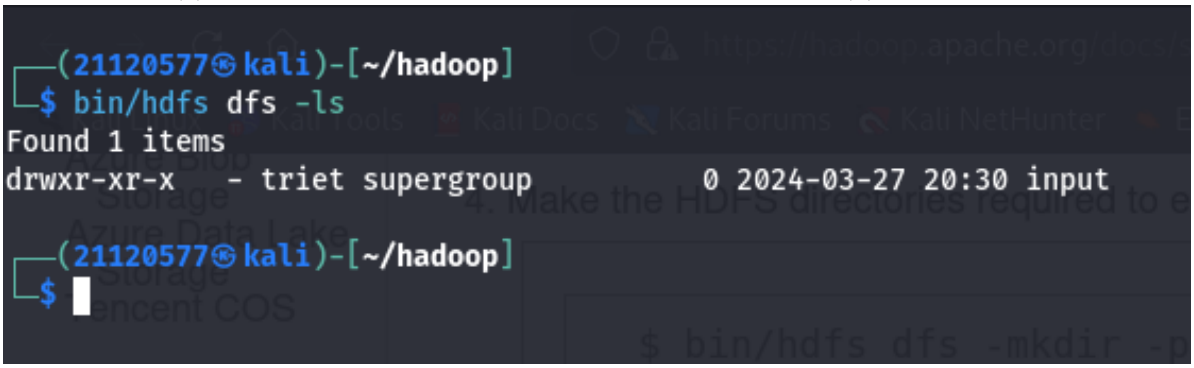

(c) Create Folders

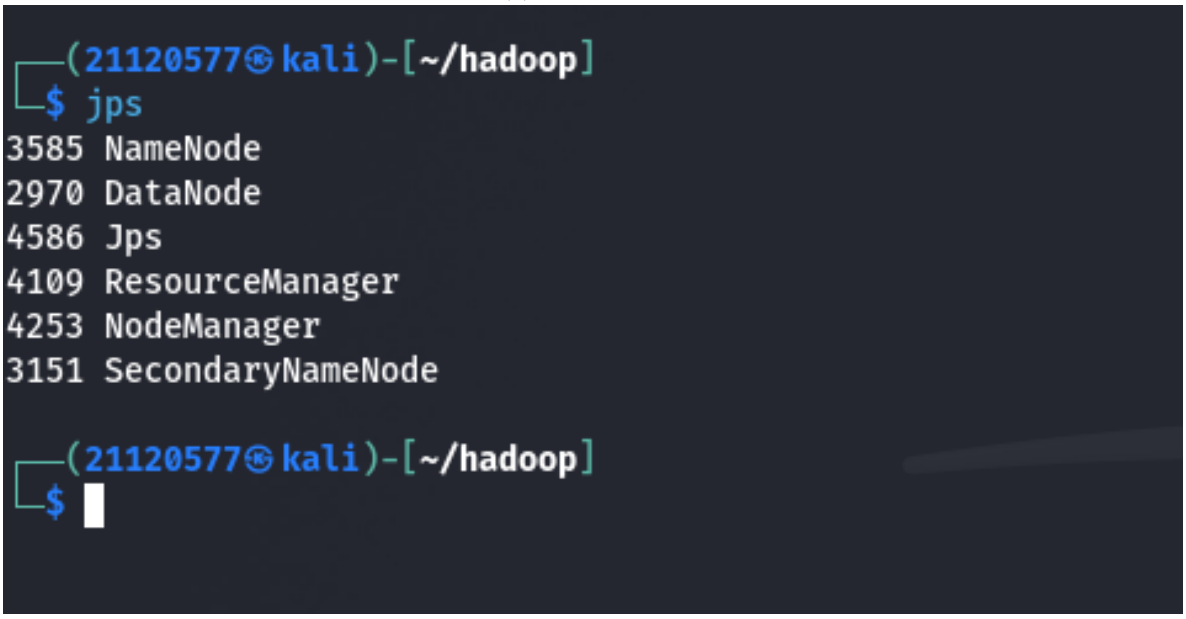

(d) JPS

Figure 9: Pseudo-Distributed Operation

2.1.4 21120583:

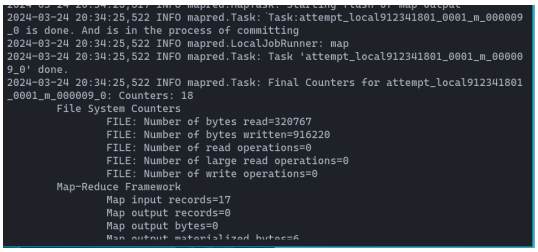

(a) hadoop-mapreduce-examples-3.3.6.jar

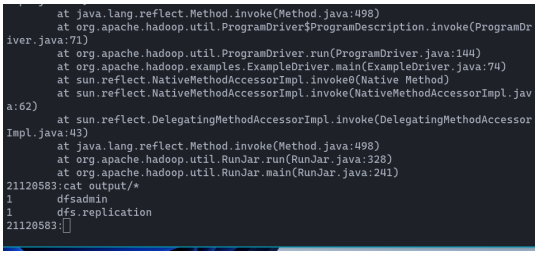

(b) List all files in output folders

Figure 10: Standalone Operation

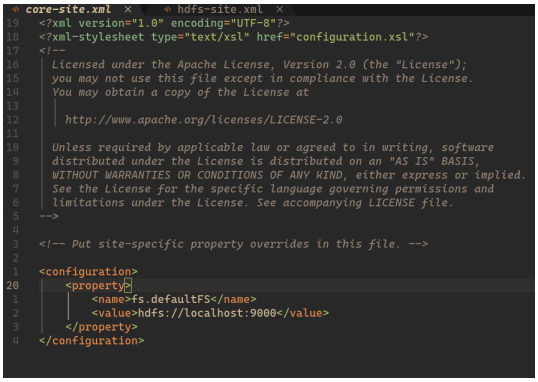

(a) etc/hadoop/core-site.xml

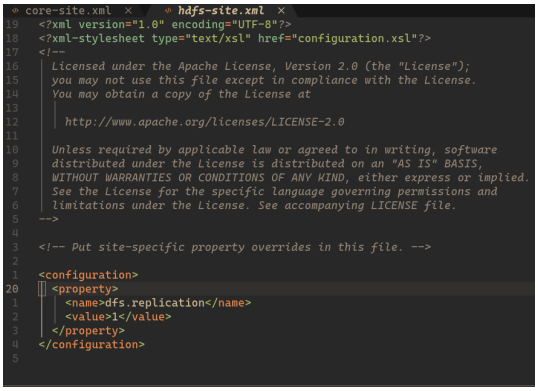

(b) etc/hadoop/hdfs-site.xml

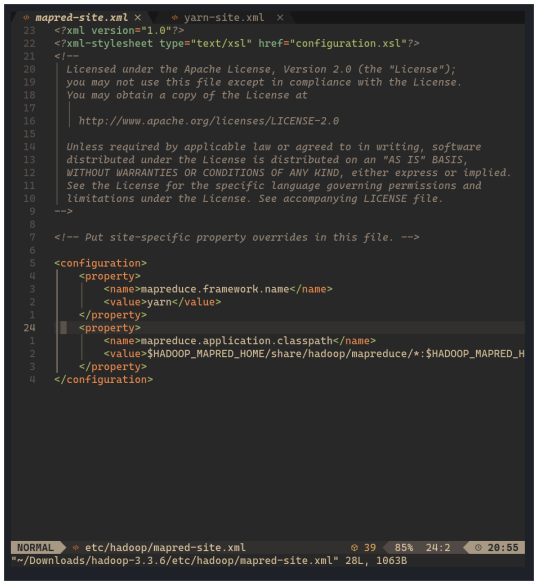

(c) etc/hadoop/mapred-site.xml

(d) etc/hadoop/yarn-site.xml

Figure 11: Hadoop configuration

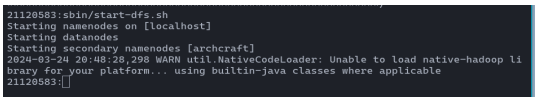

(a) Start dfs

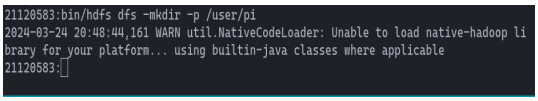

(b) Create folder

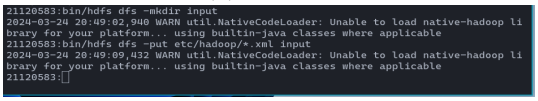

(c) Create input folder

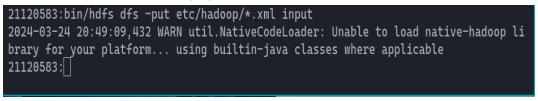

(d) Copy local .xml files to input folder

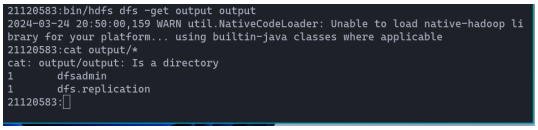

(e) Show results

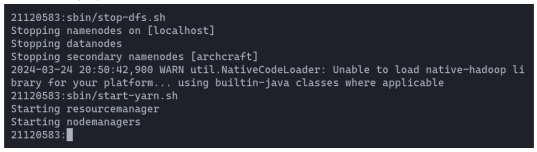

(f) Start YARN

Figure 12: Pseudo-Distributed Operation

2.2 Paper Reading

1. How do the input keys-values, the intermediate keys-values, and the output keys-values relate?

   Answer:

   MapReduce model includes 2 main operations: Map and Reduce

   - ***The Map function***: written by the user, takes an input pair and produces a set of intermediate key/value pairs.
   - ***The Reduce function***, also written by the user, accepts an intermediate key I and a set of values for that key. It merges together these values to form a possibly smaller set of values.

   The related of input keys-values, the intermediate keys-values, and the output keys-values:

   - Input keys-values represent the initial data
   - Intermediate keys-values represent the intermediate results during computation, they are essentially the output of the mapping phase and the input to the reducing phase in MapReduce
   - Output keys-values represent the final results after processing

2. How does MapReduce deal with node failures?

   Answer:

   MapReduce deal with node failures by re-executing tasks and ensuring deterministic output.There 2 main kind of failure:

   - ***Master Failure***: If the master task dies, a new copy can be started from the last checkpointed state. However, given that there is only a single master, its failure is unlikely; therefore our current implementation aborts the MapReduce computation if the master fails.
   - ***Worker Failure***: The master pings every worker periodically. If no response is received from a worker in a certacertain amount of time, the master marks the worker as failed. Any map tasks completed by the worker are reset back to their initial idle state, and therefore become eligible for scheduling on other workers. Similarly, any map task or reduce task in progress on a failed worker is also reset to idle and becomes eligible for rescheduling. Completed map tasks are re-executed on a failure because their output is stored on the local disk(s) of the failed machine and is therefore inaccessible. Completed reduce tasks do not need to be re-executed since their output is stored in a global file system.

3. What is the meaning and implication of locality? What does it use?

   Answer:

   Meaning of Locality: Conserving network bandwidth by taking advantage of the fact that the input data is stored on the local disks of the machines that make up our cluster. When running large MapReduce operations on a significant fraction of the workers in a cluster, most input data is read locally and consumes no network bandwidth.

   System tries to do the work on the computer where the data is stored. This way, it does not have to move the data over long distances, which can be slow. It's like having the toolbox (the data) right where you're working (the computer doing the task). This makes everything faster and less crowded on the network, like having a clear path in your house to walk without tripping over things. MapReduce uses this idea to make sure it can do a lot of work quickly and without causing traffic jams in the data network.

   Usage of locality in HDFS:
   Data placement policy: HDFS follows a data placement policy that tries to store replica blocks on different racks and nodes to improve locality and fault tolerance.
   Task scheduling: When scheduling tasks, HDFS tries to assign tasks to the nodes where the required data blocks are already present, taking advantage of data locality.
   Rack-aware replica placement: HDFS tries to place one replica on the same node, another on a different node in the same rack, and a third replica on a different node in a different rack, ensuring locality and fault tolerance.
   Speculative execution: If a task is running slowly due to network or hardware issues, HDFS may speculatively execute another copy of the task on a different node with better data locality, to improve overall performance.

4. Which problem is addressed by introducing a combiner function to the MapReduce model?

Answer:

> When there is significant repetition in the intermediate keys produced by each map task, and the user-specified Reduce function is commutative and associative, it costs a significant amount of time to send all intermediate keys to the Reduce function. So, a combiner function comes in handy to address this issue by partial merging of this data before it is sent over the network.
>
> The combiner function works like a mini-reducer for each map task. It takes the output of the map tasks and combines or summarizes it before it's sent over the network to the reduce tasks

2.3 Running Word Count

(a) Set env HADOOP_CLASSPATH

(b) Compile WordCount program and create file jar

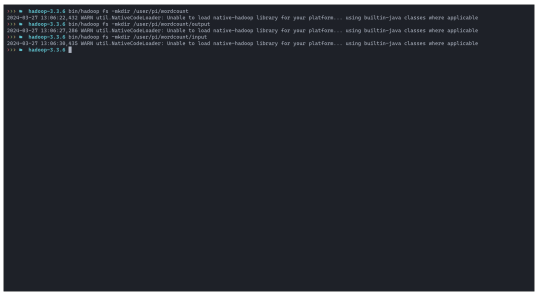

(c) Create required folders

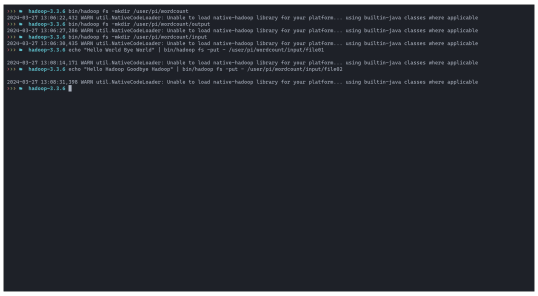

(d) Create required input files

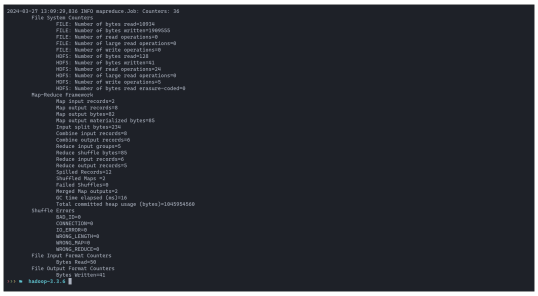

(e) Run the application

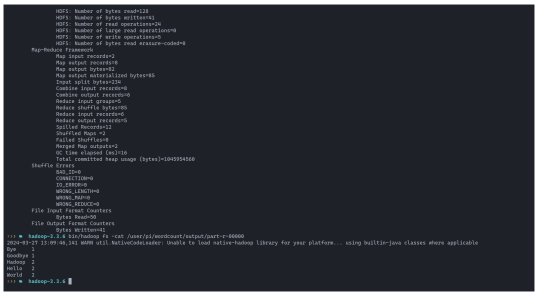

(f) output

Figure 13: Word Count

Walk through:

```
public void map(Object key, Text value, Context context)
throws IOException, InterruptedException {
  StringTokenizer itr = new StringTokenizer(value.toString());
  while (itr.hasMoreTokens()) {
    word.set(itr.nextToken());
    context.write(word, one);
  }
}
```

The Mapper operates by processing each line individually through its map method, sourced from the chosen TextInputFormat. It then breaks down each line into separate words using StringTokenizer, finally emitting key-value pairs in the format $<$ word, 1 $>$

```
public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
    IOException, InterruptedException {
  int sum = 0;
  for (IntWritable val : values) {
    sum += val.get();
  }
  result.set(sum);
  context.write(key, result);
}
```

The Reduce function computes the sum of integer values associated with a key and emits the key along with the computed sum. Keep in mind that, on this program we use Comnbiner

```
   job.setCombinerClass(IntSumReducer.class);
```

The output of each map is passed through the local combiner the function. So that, the value returned by `val.get()` may vary depending on whether it's called within the Combiner or the Reducer.

2.4 Bonus

2.4.1 Running Word Length Count

```
export JAVA_HOME="/usr/lib/jvm/java-11-openjdk-amd64"
export HDFS_NAMENODE_USER="root"
export HDFS_DATANODE_USER="root"
export HDFS_SECONDARYNAMENODE_USER="root"
export YARN_RESOURCEMANAGER_USER="root"
export YARN_NODEMANAGER_USER="root"
export PDSH_RCMD_TYPE=ssh
export PATH=${JAVA_HOME}/bin:${PATH}
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
```

(a) Set env HADOOP_CLASSPATH

```
[21120011](hadoop-3.4.0) \>
[21120011](hadoop-3.4.0) \> bin/hadoop com.sun.tools.javac.Main WordLengthCount.java
[21120011](hadoop-3.4.0) \> jar cf wc.jar WordLengthCount*.class
[21120011](hadoop-3.4.0) \>
```

(b) Compile WordCount program and create file jar

```
[21120011](hadoop-3.4.0) \>
[21120011](hadoop-3.4.0) \> bin/hadoop fs -mkdir user/root/wordlengthcount
[21120011](hadoop-3.4.0) \> bin/hadoop fs -mkdir user/root/wordlengthcount/input
[21120011](hadoop-3.4.0) \>
```

(c) Create required folders

```
[21120011](hadoop-3.4.0) \>
[21120011](hadoop-3.4.0) \>
[21120011](hadoop-3.4.0) \> echo "Hello International World I Bye National World" | bin/hadoop fs -put - user/root/wordlengthcount/input/file01
[21120011](hadoop-3.4.0) \> echo "I am Backtracking Master" | bin/hadoop fs -put - user/root/wordlengthcount/input/file02
[21120011](hadoop-3.4.0) \> echo "Haha Hihi 10 overall Big Data" | bin/hadoop fs -put user/root/wordlengthcount/input/file03
put: `user/root/wordlengthcount/input/file03': No such file or directory
[21120011](hadoop-3.4.0) \> echo "Haha Hihi 10 overall Big Data" | bin/hadoop fs -put - user/root/wordlengthcount/input/file03
[21120011](hadoop-3.4.0) \>
```

(d) Create required input files

```
                Input split bytes=408
                Combine input records=17
                Combine output records=10
                Reduce input groups=4
                Reduce shuffle bytes=135
                Reduce input records=10
                Reduce output records=4
                Spilled Records=20
                Shuffled Maps =3
                Failed Shuffles=0
                Merged Map outputs=3
                GC time elapsed (ms)=231
                CPU time spent (ms)=6230
                Physical memory (bytes) snapshot=1147174912
                Virtual memory (bytes) snapshot=11012554752
                Total committed heap usage (bytes)=829423616
                Peak Map Physical memory (bytes)=317546496
                Peak Map Virtual memory (bytes)=2751893504
                Peak Reduce Physical memory (bytes)=209846272
                Peak Reduce Virtual memory (bytes)=2760695808
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=102
        File Output Format Counters
                Bytes Written=30
[21120011](hadoop-3.4.0) \>
```

Figure 15: Run the application

```
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=102
        File Output Format Counters
                Bytes Written=30
[21120011](hadoop-3.4.0) \> bin/hadoop fs -ls user/root/wordlengthcount/output
Found 2 items
-rw-r--r--   1 root supergroup          0 2024-03-30 19:19 user/root/wordlengthcount/output/_SUCCESS
-rw-r--r--   1 root supergroup         30 2024-03-30 19:19 user/root/wordlengthcount/output/part-r-00000
[21120011](hadoop-3.4.0) \> bin/hadoop fs -cat user/root/wordlengthcount/output/part-r-00000
Big     2
Medium  6
Small   7
Tiny    2
[21120011](hadoop-3.4.0) \>
```

Figure 16: output

Walk through:

```
public void map(Object key, Text value, Context context
) throws IOException, InterruptedException {
    StringTokenizer itr = new StringTokenizer(value.toString());
    while (itr.hasMoreTokens()) {
        String token = itr.nextToken();
        String type = classifyWordLength(token);
        word.set(type);
        context.write(word, one);
    }
}

private String classifyWordLength(String token) {
    if(token.length() == 1) return "Tiny";
    else if (token.length() >= 2 && token.length() <= 4) {
        return "Small";
    } else if (token.length() >= 5 && token.length() <= 9) {
        return "Medium";
    }
    return "Big";
}
```

We employs a custom Mapper class named **TokenizerMapper**, which extends Hadoop's **Mapper** class. This class is responsible for reading input text data, tokenizing the text into individual words, and categorizing each word based on its length. For each word, the **map** method emits a key-value pair, with the key being one of the predefined categories ("Tiny", "Small", "Medium", "Big") and the value being an **'IntWritable'** of 1, representing a single occurrence of a word in that category.

```
public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
    IOException, InterruptedException {
  int sum = 0;
  for (IntWritable val : values) {
    sum += val.get();
  }
  result.set(sum);
  context.write(key, result);
}
```

The **IntSumReducer** class extends Hadoop's **Reducer** class and serves to aggregate the counts of words falling into each category. It takes the intermediate key-value pairs produced by the Mapper as input and sums up the values for each category, resulting in a total count of words for each word length category. This aggregated result is then written to the output, providing a comprehensive count of words per category.

2.4.2 Setting up Fully Distributed Mode

a) Hadoop Cluster Setup in Non-Secure Mode:

1. Create user named "hadoop" on all machines
2. Make sure our machines can run ssh localhost.
3. Connect ssh to workers without password.
4. Edit configuration files on Master and Slave node:
   (a) Master:

(a) core-site.xml

(b) hdfs-site.xml

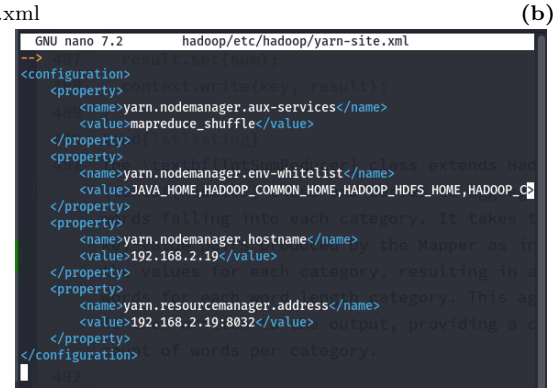

(c) yarn-site.xml

Figure 17: Namenode configuration

(b) Slave:

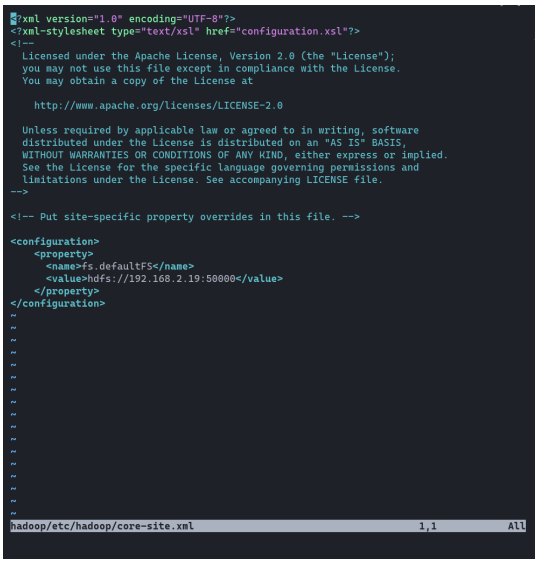

(a) core-site.xml

(b) hdfs-site.xml

(c) yarn-site.xml

Figure 18: Worker-node configuration

5. Format namenode again
6. Start hadoop on Master node

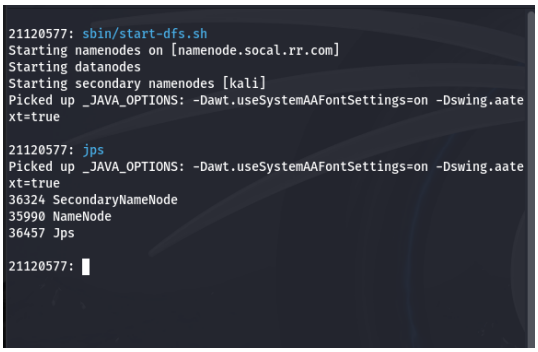

(a) Start DFS

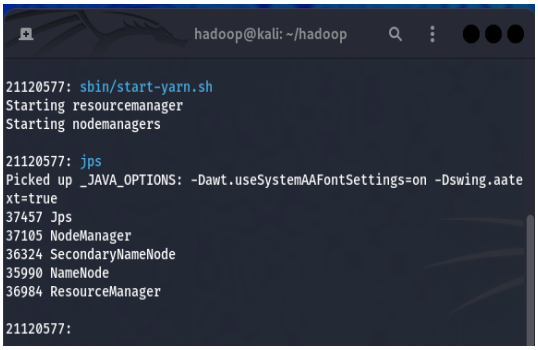

(b) Start Yarn

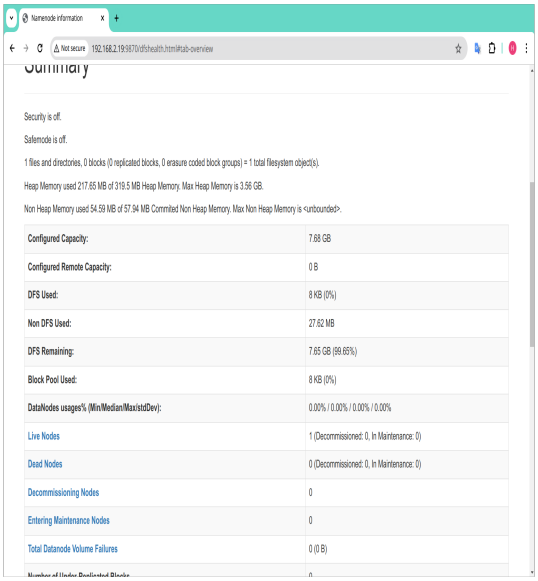

(c) Check connection in UI

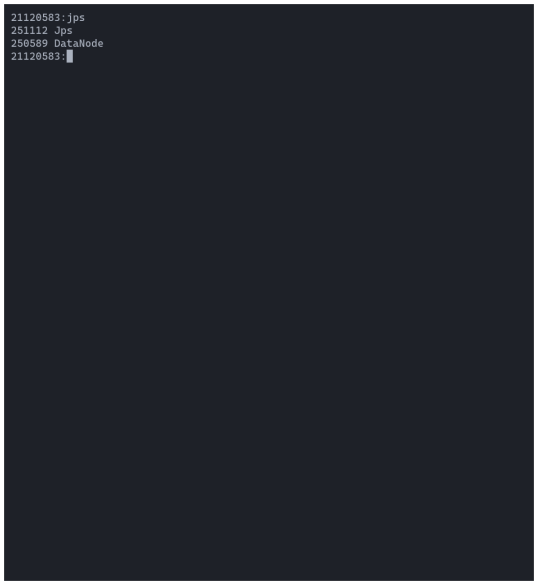

(d) Check datanode start at worker machine

Figure 19: Verify namenode start on slave successfully

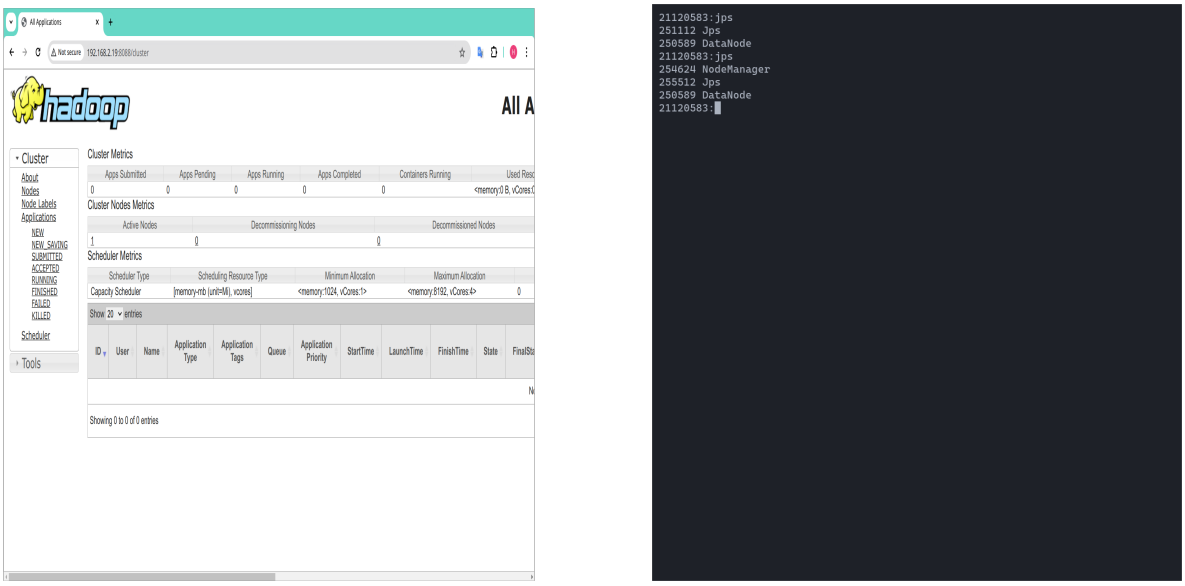

Figure 20: Verify YARN start on slave successfully

Note:

- To avoid unexpected behaviours, we should create the same user name for all machines.
- All machine must connect the same local network.
- Installing the same Hadoop version on each machine is required

b) Research about Security in Hadoop Set-up

1. Is your Hadoop secure? Give a short explanation if your answer is yes. Otherwise, give some examples of risks to your system.
   Answer:

The Hadoop system would not be considered secure. Here's some example of risks to my system:

User to Service Authentication:
+ No User Authentication on NameNode or JobTracker: If there is no mechanism to authenticate users, then any client can potentially claim to be any user, which is a significant security risk.
+ No User Authorization on DataNode: Prior to version 0.21, users could read/write any block, leading to the possibility of unauthorized data access and potential data leakage.
+ No User Authorization on JobTracker: Without proper authorization, users could modify or kill jobs that do not belong to them, affecting the data processing integrity and potentially causing a denial of service.

Service to Service Authentication:
+ No Authentication of DataNodes and TaskTrackers: Without authentication, there is a risk of unauthorized nodes joining the cluster and interfering with its operation or stealing data.

No Encryption on Wire or Disk:
+ Data transmitted across the network (on the wire) or stored on disk without encryption can be intercepted or accessed by unauthorized entities, leading to data breaches.

2. From your perspective, which method is better when securing your HDFS: authentication, authorization, or encryption? Give an explanation about your choices.
Answer:

Authentication ensures that only identified and verified users can access the system. It's the first line of defense, preventing unauthorized access.

Authorization goes one step further by defining access rights and permissions for authenticated users. This ensures that users can only perform actions permitted by their role, which is essential for enforcing the principle of least privilege.

Encryption protects data at rest and in transit, ensuring confidentiality and integrity. Even if an attacker circumvents authentication and authorization mechanisms, encryption can prevent them from reading or tampering with the data.

Each of these methods serves a different purpose:
- Authentication protects against unauthorized entry into the system.
- Authorization ensures that users can't perform actions they're not permitted to do.
- Encryption keeps the data secure even if the system is breached.

In the context of a student project, authentication would be the most suitable method to prioritize when securing Hadoop Distributed File System (HDFS).
Authentication ensures that only authorized individuals can access the HDFS by verifying their identity. By implementing strong authentication mechanisms, such as username/-password authentication or integration with external authentication systems like LDAP, the project can enforce access controls and prevent unauthorized users from gaining access to the HDFS.

3 Reflection

3.1 Journey to the Deadline:

Our work is divided into 4 main phase:

1. *Each member implements Standalone Operation and Pseudo-Distributed Operation of Hadoop.*
2. *Reading Paper and Answer question about MapReduce model*
3. *Running example of Word Count program and implementing MapReduce programs: Word Length Count*
4. *Implementing Fully-connected Distributed Operation on 2 physical devices and researching about Security in Hadoop Set-up*

Phase	Time-consumption	Bugs
1	~6 days	+ ssh: connection denied port 22 + java version is not compatible with hadoop version + cannot create folder in hdfs
2	~2 days	
3	~2 days	
4	~3 days	+ cannot start datanode on worker device + cannot set up the datanode workspace

3.2 Overcoming Challenges

1. SSH problem: we have researched the problem and found that it can be solve by restarting ssh service.

```
sudo service ssh restart
```

```
ssh localhost
```

2. java version is not compatible with hadoop version (3.4.0). We found that cannot run hdfs with jdk-21 so we solved by installing jdk-11.
3. Cannot create folder on hdfs. The problem is default root of hdfs which start `/user/username` so we need create that folder first of add / before any command.
4. Cannot start datanode on worker device. After configurating, we did not connect ssh between master device and worker device that why master node cannot connect to worker node. Another problem is the permission of hadoop folder is root. We solved by run commands change owner of hadoop folder

```
sudo chown hadoop:hadoop -R /home/hadoop
```

```
chmod 700 /home/hadoop
```

5. Cannot set up the datanode workspace. The default workspace of datanode is `/home/username/hadoop/bin/hdfs`, so when we start datanode from master we cannot find `/home/master-username/hadoop/bin/hdfs` on worker device and we try to configure `hdfs-site.xml` but it's still not working. We have to create the user named hadoop on both master and worker device to solved this problem.

```
Starting datanodes
pi@datanode1: zsh:1: no such file or directory: /home/triet/hadoop/bin
/hdfs
pdsh@kali: pi@datanode1: ssh exited with exit code 127
```

Figure 21: No such file error

3.3 Lesson Learned

After setting up a Hadoop cluster and running a MapReduce program, there are several valuable lessons that can be learned. Firstly, the process of setting up the cluster itself provides insights into distributed systems and infrastructure management. Understanding the configuration and interplay of various components such as NameNode, DataNode, ResourceManager, and NodeManager enhances knowledge of cluster architecture.

Running a MapReduce program offers practical experience in processing large-scale data in a distributed environment. It highlights the importance of data partitioning, task scheduling, and fault tolerance mechanisms. Dealing with input/output formats, understanding the Map and Reduce functions, and optimizing data locality contribute to a deeper comprehension of parallel processing and data manipulation.

Moreover, resolving the problems encountered during the setup of the Hadoop cluster can significantly improve our team's Linux command skills. The process of configuring and troubleshooting the cluster often involves working with various Linux commands and utilities. Setting up a Hadoop cluster typically requires tasks such as installing and configuring the operating system, managing user permissions, modifying network settings, and installing and configuring software packages. These tasks often involve executing Linux commands to navigate the file system, modify configuration files, manage services, and troubleshoot issues.

4 Reference

- Hadoop - Multi-Node Cluster
- Hadoop Multi-node Cluster Installation
- Viblo - Set-up multi node cluster hadoop
- Safe Mode in Hadoop
- Apache Hadoop 3.3.6 – Hadoop Cluster Setup
- Apache Hadoop 3.3.6 – Hadoop: Setting up a Single Node Cluster
- Apache Hadoop 3.3.6 – Hadoop in Secure Mode
- Apache Hadoop 3.3.6 – MapReduce Tutorial